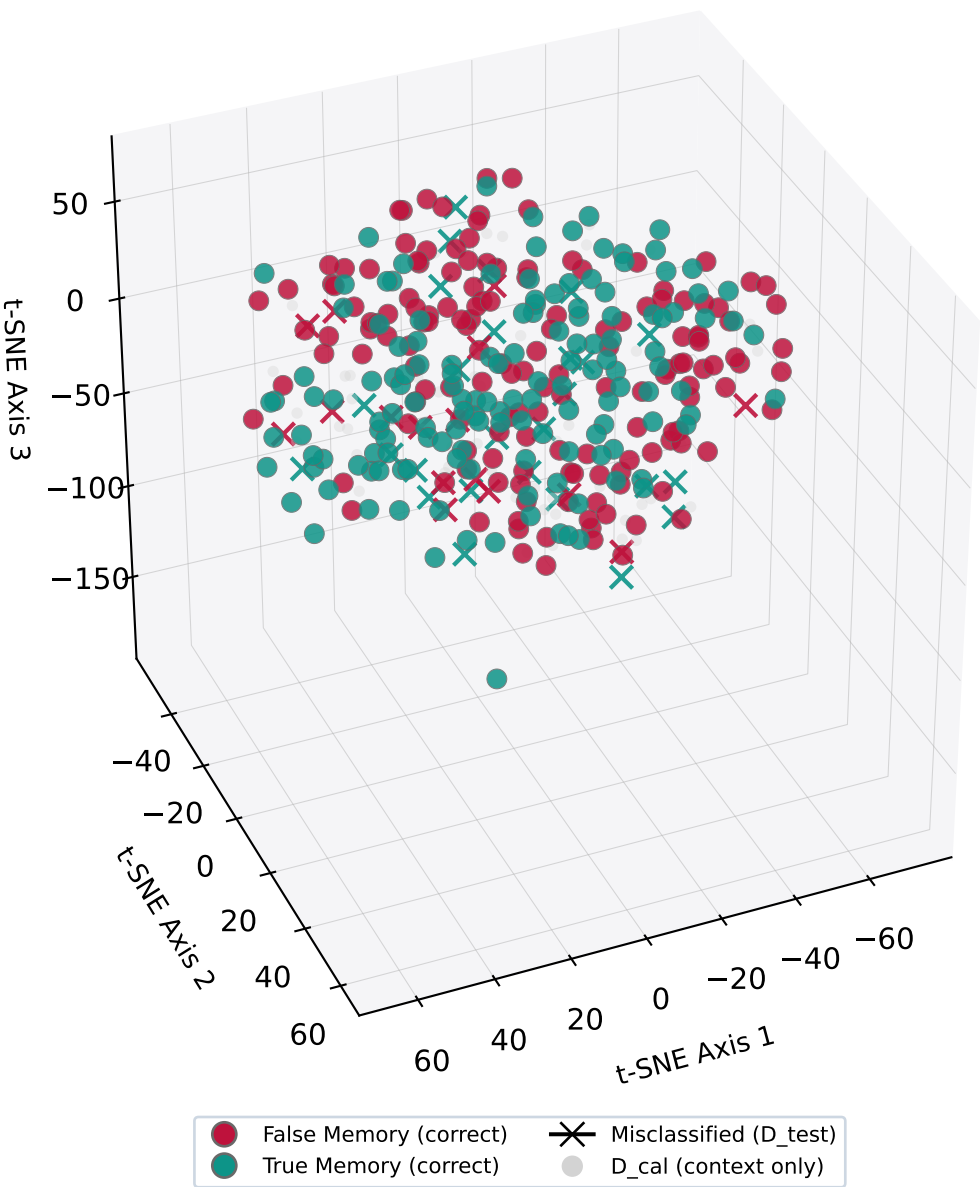
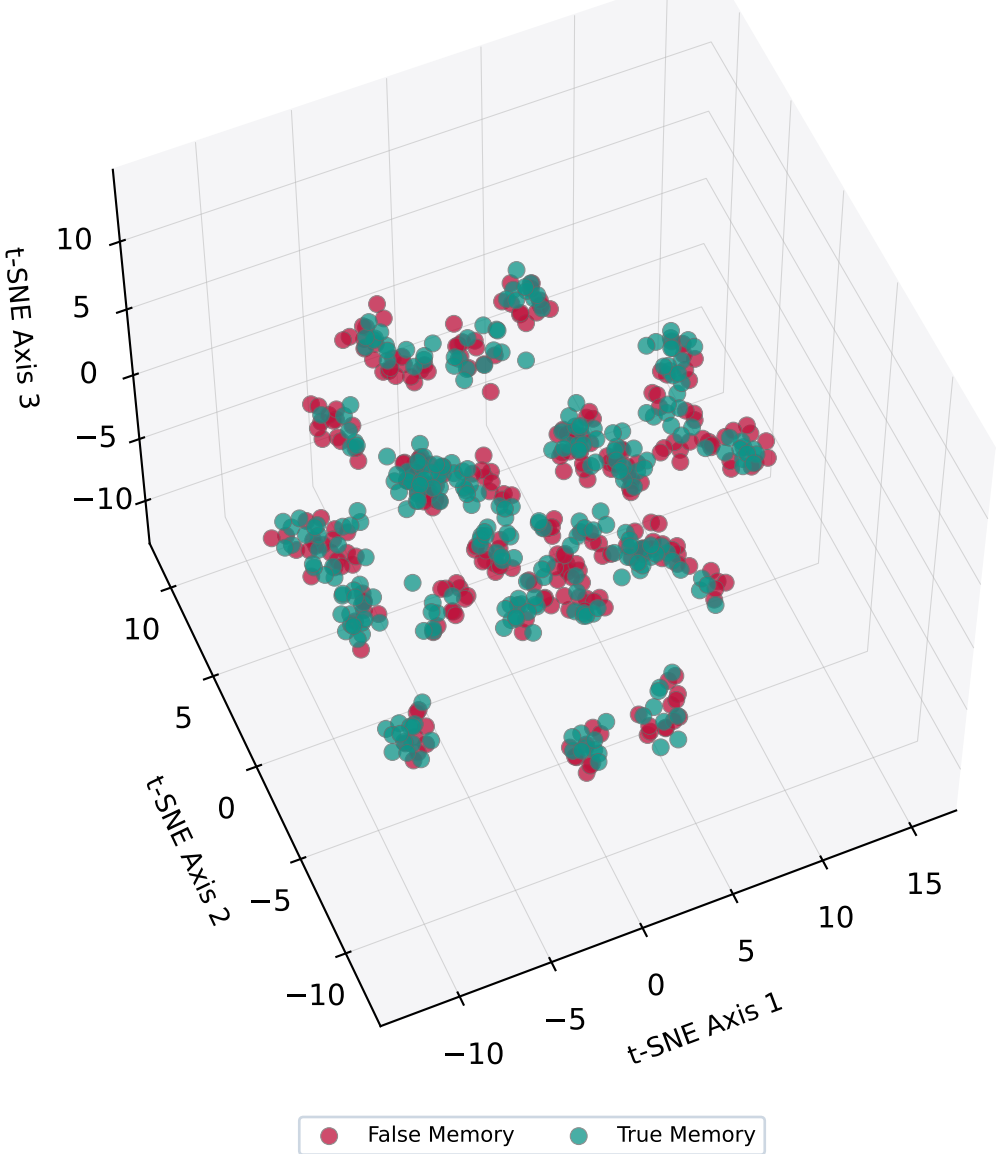


Figure 3: Best-Case Single-Subject Separability vs. Population-Level Subject-Islands Effect

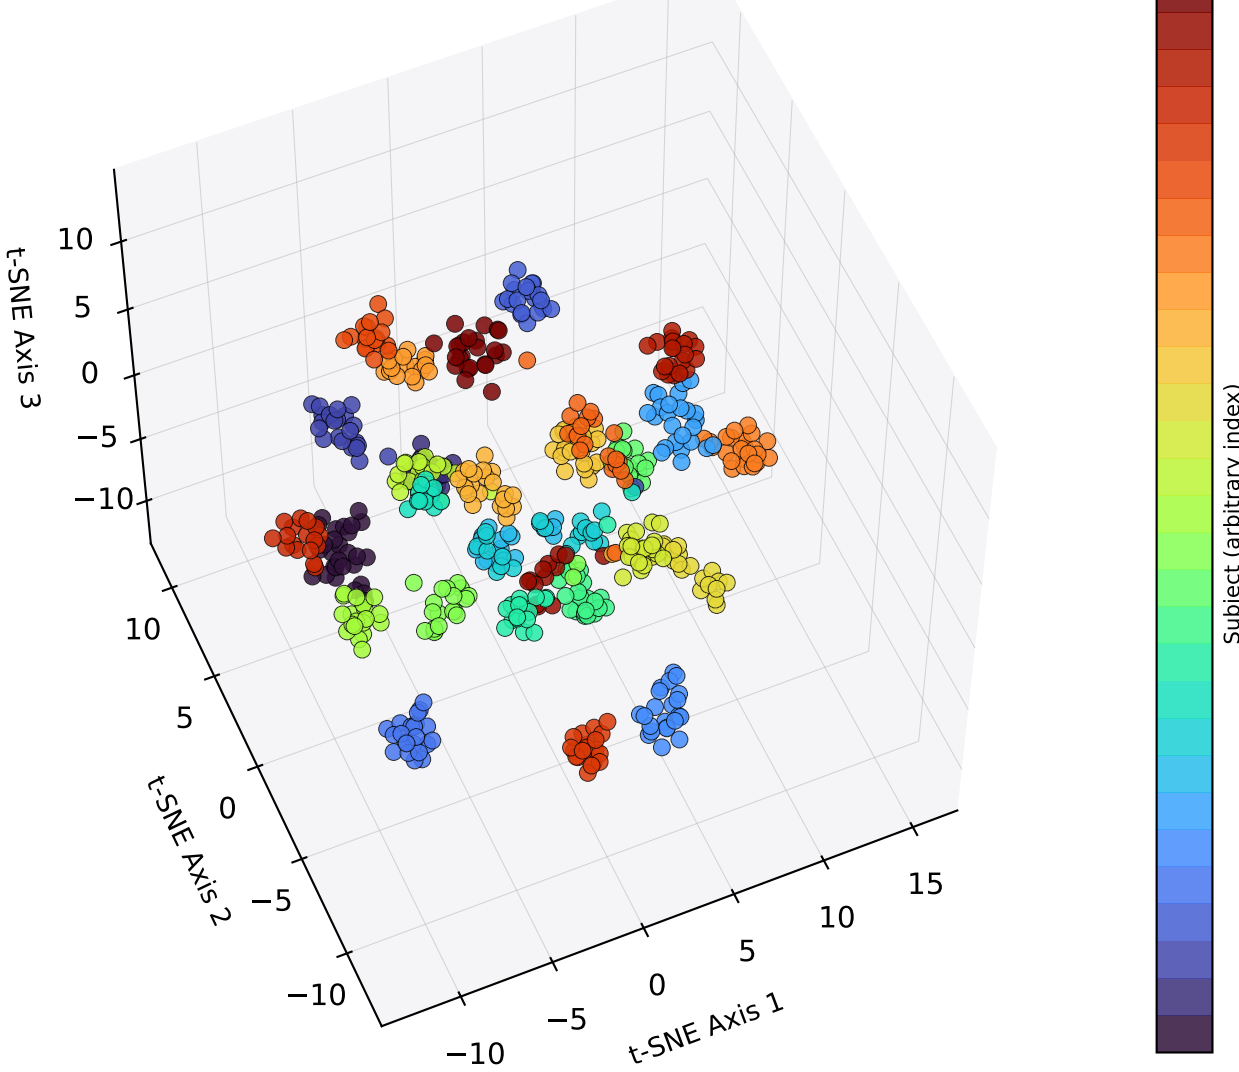
A. Best-Responder Example: sub-11
(rank 1/29, held-out acc=87.6%; group mean=67.8%±8.7%)



B. Multi-Subject (N=600), Colored by Class
(label-informed axes; centroid gap=1.22)



C. Same Coordinates, Colored by Subject ID
(29 subjects → geometry tracks identity, not class)



A: Unsupervised PCA axes (fit on 28 subjects only) + t-SNE layout; markers show real held-out prediction correctness.
B/C: Label-informed LDA+PCA → t-SNE axes on pooled pre-calibration features, demonstrating geometry tracks subject identity prior to calibration.